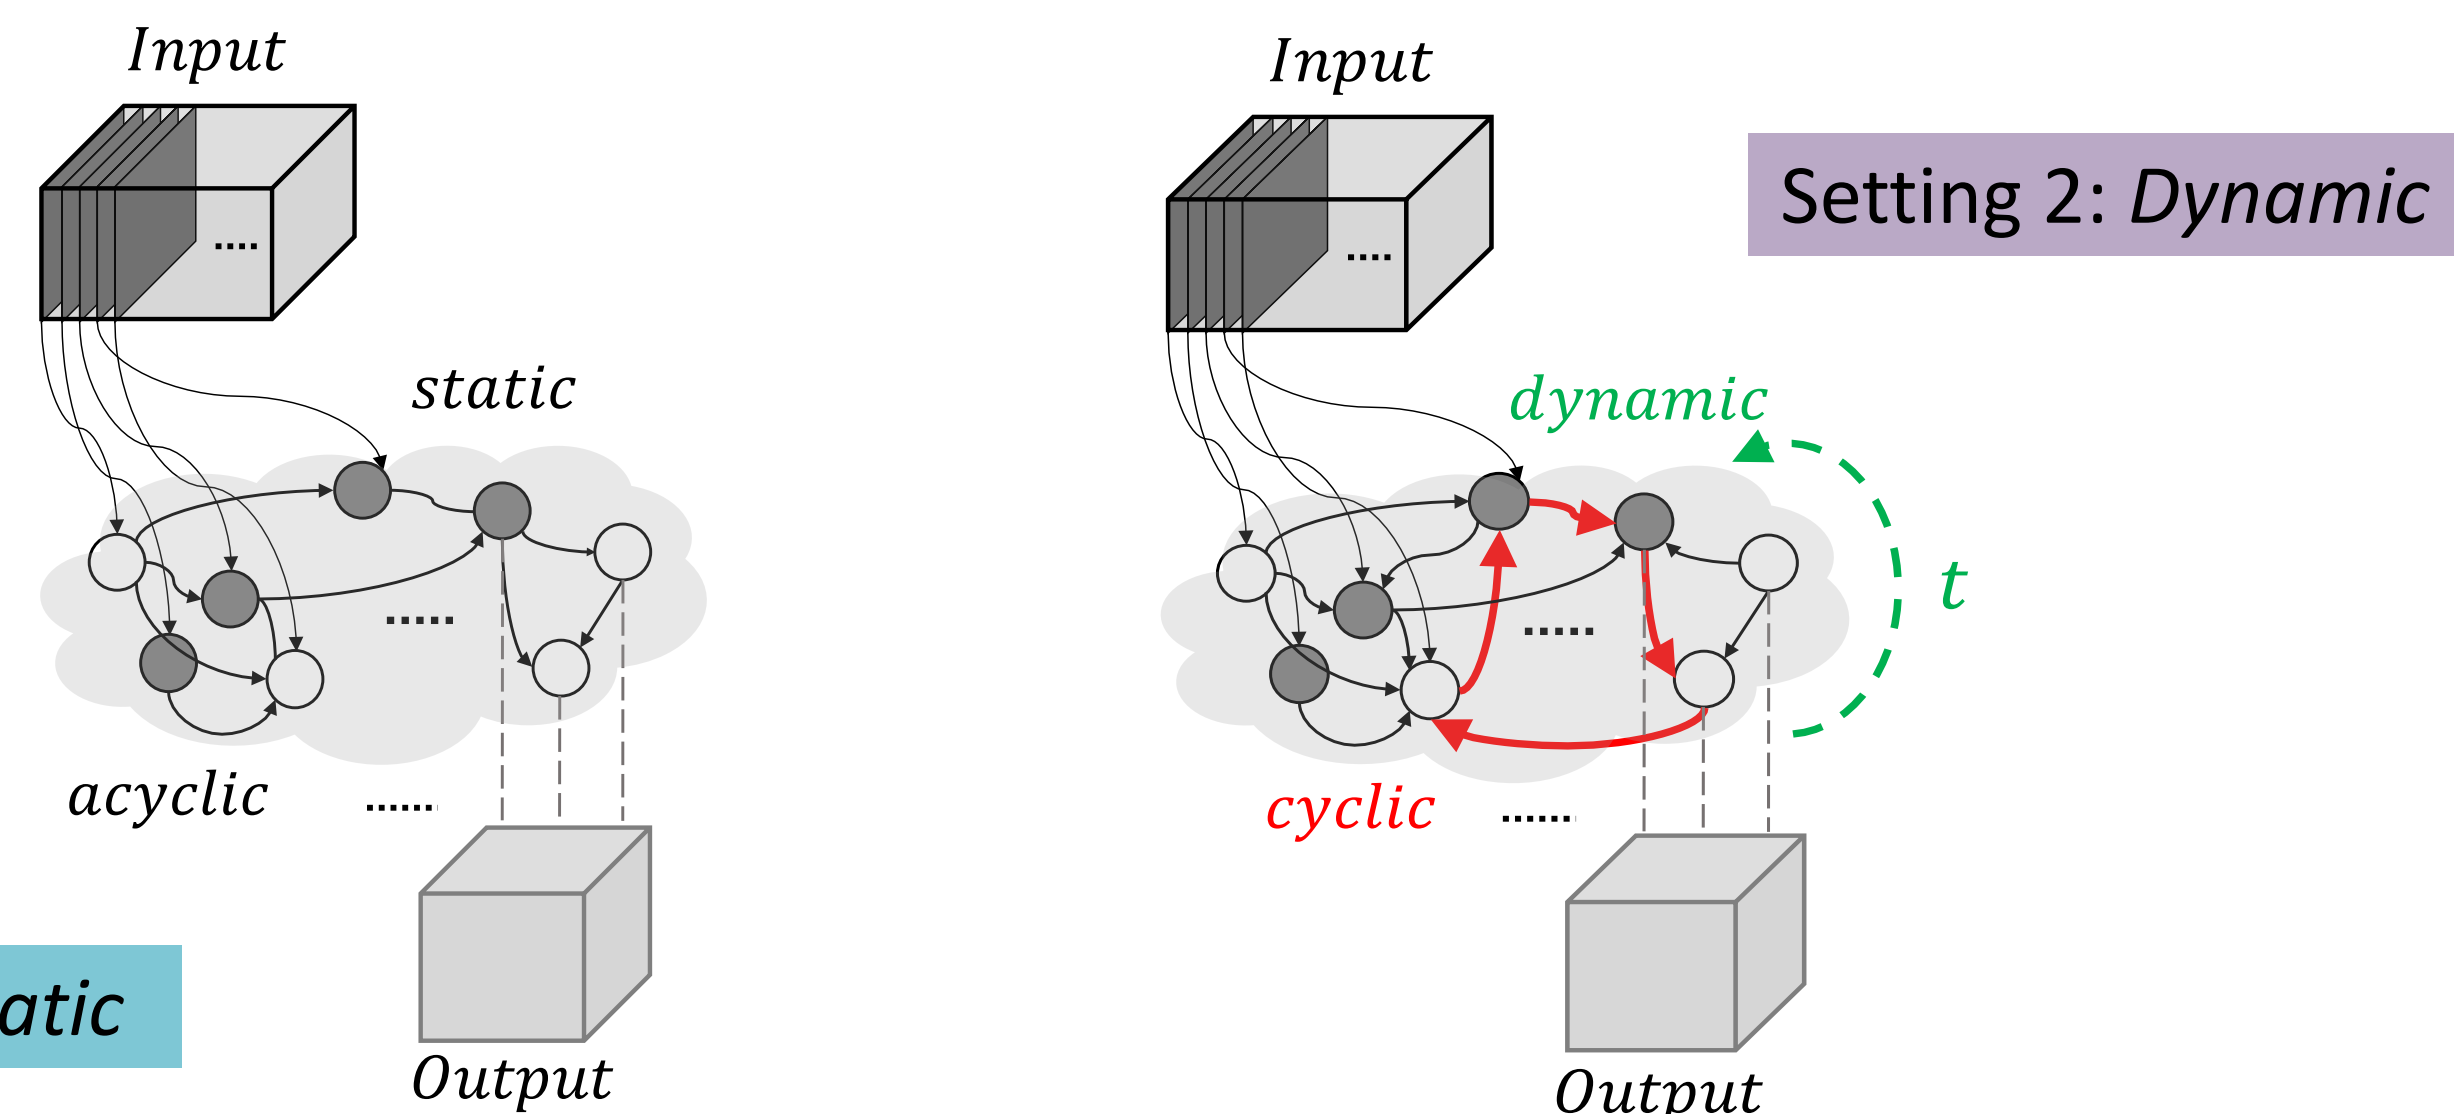


MOTIVATION

- Traditionally, the connectivity patterns of neural networks are manually defined or largely constrained (even with methods of Neural Architecture Search (NAS) [1]).
- We allow for a much **larger space of possible networks** by relaxing the typical notion of layers and enabling channels to form connections independent of each other.
- The **wiring of our network is not fixed** during training – as we learn the network parameters we also learn the connectivity.




CONTRIBUTIONS

- We present an algorithm, **Discovering Neural Wirings (DNW)**, to efficiently learn the connectivity of a neural network.
- By learning the connectivity of MoileNetV1 (x0.25) [2] we **boost the accuracy on ImageNet by 10%**.
- We demonstrate that DNW can also be used to **effectively train sparse neural networks in a single training run**: We only ever train with 10% of weights (in the forward pass) and only lose 2.5% accuracy on ImageNet compared to our dense baseline.

KEY TAKEAWAYS

- It is possible to realize the benefits of overparameterization during training, even when the resulting model is sparse.
- As NAS becomes more fine grained, finding a good architecture is akin to finding a sparse subnetwork of the complete graph.

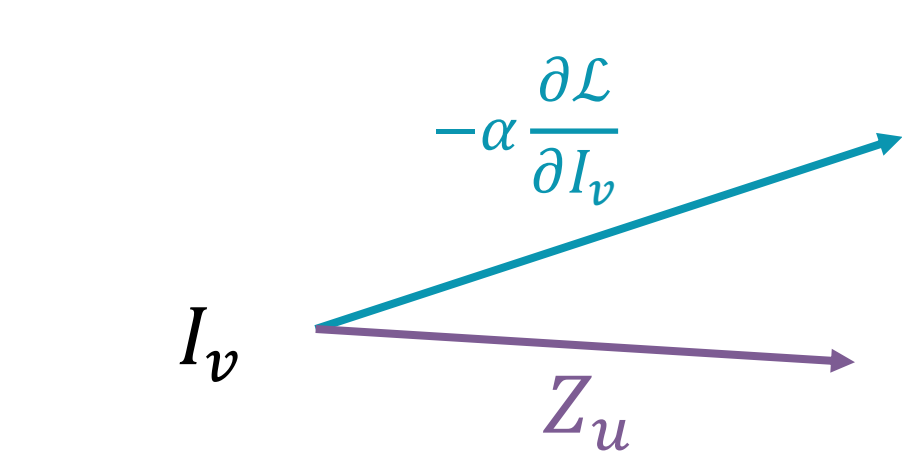


- We consider the case where the input and output of each node is a two-dimensional matrix, commonly referred to as a channel.
- The state of a fixed set of output nodes are taken to be the output of the network.
- We let f_{θ_u} at each non-output node be a batch-norm, (2 parameters), ReLU, 3x3 convolution (9 parameters) triplet.
- The goal is to learn a good graph with exactly k edges.

Intuition

- If the gradient is pushing I_v in a direction which aligns with Z_u , then we strengthen the magnitude of the weight w_{uv} .
- If this alignment happens consistently then $|w_{uv}|$ will be eventually be strong enough to enter the edge set \mathcal{E} .
- If (u, v) enters \mathcal{E} , another edge will be removed. We show that when this swapping does occur, it is beneficial (i.e. it decreases the loss on the mini-batch under certain conditions).
- If the edge is already in the graph, the update rule is no different than standard SGD.

Update Rule

$$w_{uv} \leftarrow w_{uv} + \left\langle Z_u, -\alpha \frac{\partial \mathcal{L}}{\partial I_v} \right\rangle$$


Results

By learning the connectivity of MobileNetV1 [3], we boost the ImageNet accuracy by ~10% at ~41M FLOPS.

Model	Params	FLOPs	Accuracy
MobileNetV1 ($\times 0.25$)	0.5M	41M	50.6%
X-4 MobileNetV1	—	> 50M	54.0%
MobileNetV2 ($\times 0.15$)*	—	39M	44.9%
MobileNetV2 ($\times 0.4$)**	—	43M	56.6%
DenseNet ($\times 0.5$)*	—	42M	41.1%
Xception ($\times 0.5$)*	—	40M	55.1%
ShuffleNetV1 ($\times 0.5, g = 3$)	—	38M	56.8%
ShuffleNetV2 ($\times 0.5$)	1.4M	41M	60.3%
MobileNetV1-Random Graph($\times 0.225$)	1.2M	55.7M	53.3%
MobileNetV1-DNW-Small ($\times 0.15$)	0.24M	22.1M	50.3%
MobileNetV1-DNW-Small ($\times 0.225$)	0.4M	41.2M	59.9%
MobileNetV1-DNW ($\times 0.225$)	1.1M	42.1M	60.9%
MnasNet-search1	1.9M	65M	64.9%
MobileNetV1-DNW($\times 0.3$)	1.3M	66.7M	65.0%
MobileNetV1 ($\times 0.5$)	1.3M	149M	63.7%
MobileNetV2 ($\times 0.6$)*	—	141M	66.6%
MobileNetV2 ($\times 0.75$)*	—	145M	67.9%
DenseNet ($\times 1$)*	—	142M	54.8%
Xception ($\times 1$)*	—	145M	65.9%
ShuffleNetV1 ($\times 1, g = 3$)	—	140M	67.4%
ShuffleNetV2 ($\times 1$)	2.3M	146M	69.4%
MobileNetV1-Random Graph($\times 0.49$)	1.8M	170M	64.1%
MobileNetV1-DNW($\times 0.49$)	1.8M	154M	70.4%

Comparing the **Static** vs. **Dynamic** Setting with a tiny (41k parameter) classifier on CIFAR-10.

Model	Accuracy
Static (Random Graph)	76.1 \pm 0.5%
Static (DNW)	80.9 \pm 0.6%
Discrete Time (Random Graph)	77.3 \pm 0.7%
Discrete Time (DNW)	82.3 \pm 0.6%
Continuous (Random Graph)	78.5 \pm 1.2%
Continuous (DNW)	83.1 \pm 0.3%


Contrasting with other methods of Discovering Wirings with a fixed structure and number of edges on CIFAR-10.

Model	Accuracy
MobileNetV1 ($\times 0.25$)	86.3 \pm 0.2%
MobileNetV1-Random Graph($\times 0.225$)	87.2 \pm 0.1%
No Update Rule	86.7 \pm 0.5%
L1 + Anneal	84.3 \pm 0.6%
Targeted Dropout $\rho = 0.95$	89.2 \pm 0.4%
Lottery Ticket (one-shot)	87.9 \pm 0.3%
MobileNetV1-DNW ($\times 0.225$)	89.7 \pm 0.2%

Sparse Neural Network Learning

Method	Weights (%)	Top-1 Accuracy	Top-5 Accuracy
Sparse Networks from Scratch	10%	72.9%	91.5%
Ours - All Layers Sparse	10%	74.0%	92.0%
Ours - First Layer Dense	10%	75.0%	92.5%
Sparse Networks from Scratch	20%	74.9%	92.5%
Ours - All Layers Sparse	20%	76.2%	93.0%
Ours - First Layer Dense	20%	76.6%	93.4%
Sparse Networks from Scratch	30%	75.9%	92.9%
Ours - All Layers Sparse	30%	76.9%	93.4%
Ours - First Layer Dense	30%	77.1%	93.5%
Sparse Networks from Scratch	100%	77.0%	93.5%
Ours - Dense Baseline	100%	77.5%	93.7%

We apply our algorithm to the task of training a sparse ResNet-50. The sparsity is maintained throughout training, as motivated by Sparse Networks From Scratch [4].



[1] Barret Zoph and Quoc V. Le. **Neural Architecture Search with Reinforcement Learning**. In *ICLR*, 2016.

[2] Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt and David Duvenaud. **Neural Ordinary Differential Equations**. In *NeurIPS*, 2018.

[3] Andrew G. Howard et al. **MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications**. In *ArXiv*, 2017.

[4] Tim Dettmers and Luke Zettlemoyer. **Sparse Networks from Scratch: Faster Training without Losing Performance**. In *ArXiv*, 2019.

